

Lecture 17



Presentation Details

Presentation Requirements (1 of 2)

- Presentation order: Decreasing order of team number
- Code cutoff is *this* Sunday at midnight
 - Have to ask permission to commit code afterwards, if your presentation depends on it
 - Spend time rehearsing instead



Presentation Requirements (1 of 2)

- Readme files are now required inside Sunday's project submission.
 - Text summary of:
 - Non-obvious things implemented
 - Difficulties
 - Advanced topics used
 - Teammates' contributions



Presentation Requirements (2 of 2)

- Teammate surveys due Friday of presentations
 - Submit a short text file assigning a score to each teammate, including a short comment about why



Presentation Requirements (2 of 2)

- Plan how to set up for the projector quickly
- Find adapters for your computer!!!!
 - Otherwise there's only my computer, navigated to your URL
 - My computer is slow
 - Your GitHub pages must be flipped on, otherwise even that won't work
 - Remember that enabling GitHub Pages is required
- If possible swing by this room
 - While it's still open, when no one is in here
 - Try plugging in.



Grading

Announcements

- Grading is late except for midterm
 - Don't panic about the impression that gives
 - Expect good grades in the class
 - ***Project grades will average high***
 - Final grades too
 - You earned it





Course Evaluations (Required!)

Announcement

- 2 free percentage points in the class from filling out a course evaluation
- (20 points on our 1000 point scale)
- This is another attempt to hit the 1100 points total required for A+ in the grading scale (that I promised would be used)
- Important: Do not forget to fill one out!





Final Exam Info and Tips

Final Exam Format

- 40-50 multiple choice questions
- Scantron again
- Cumulative coverage (early stuff and late stuff)
- 11:30-2:30 Friday June 14
 - Young Hall CS76



Final Exam Coverage

- The midterm's harder questions: Transformation ordering, moving between bases
- What's one sequence of matrices that can be used to connect two squares by a corner?
- Implicit, Explicit, and Parametric -- How do you convert between any pair of these equation types?
- Simple parametric shapes $f(s,t)$
- Hermite spline curves; G1 vs C1 continuity



Final Exam Coverage

- What sequence of matrices reflects point P around an arbitrary plane, given the plane's equation?
- What sequence of matrices change a point's basis from any pair of arrows A to a pair B ? (Be careful: Try a simple case first where A is just a single translation away from being B)
- What were the formulas used for ray tracing?
- What are the dependent variables when trying to find where a line's vanishing point is?



Rule of Thumb

Transforming a point P :

Transformations: T_1, T_2, T_3

Matrix: $\mathbf{M} = \mathbf{M}_3 \mathbf{M}_2 \mathbf{M}_1$

Point is transformed by $\mathbf{M}P$

Each transformation happens with respect to the **same** coordinate system

Transforming a coordinate system:

Transformations: T_1, T_2, T_3 (not generally the same as the ones above)

Matrix: $\mathbf{M} = \mathbf{M}_1 \mathbf{M}_2 \mathbf{M}_3$

A point has coordinates $\mathbf{M}P$ in the original coordinate system

Each transformation happens with respect to **previous** coordinate system

Rule of Thumb

To find the transformation matrix that transforms P from C_A coordinates to C_B coordinates, we find a sequence of transformations that align C_B to C_A , accumulating matrices from left to right

Explanation of This Rule

If we think coordinate systems, **M** takes C_A from the left and produces C_B on the right:


$$C_A \mathbf{M}_B = C_B$$

After this transformation we “talk” in C_B coordinates (right side).

If we think points, then we go the other way; **M** takes P_B on the right and produces the P_A coordinates on the left:


$$P_A = \mathbf{M}_B P_B$$

Transformation **M**: \mathbf{M}_B



Explanation of This Rule

Transformation \mathbf{M} : ${}_A\mathbf{M}_B$



Consider this simple example, where to produce C_B we translate C_A by +1 along the x axis:

$$P_A = (2,1) \quad P_B = (1,1)$$

If we move C_A by +1 in x to transform it into C_B then the x coordinate of P with respect to the new system is reduced by 1 (C_B is closer to P than C_A by 1).

So, if we want to transform the coordinates of P from C_B to C_A we need to add 1 in x. Exactly what we need to do to transform C_A to C_B .

Final Exam Coverage

- That's all my tips
- Study the slides! They're the main source of questions.





Prevent slowdown in your project: Performance profiling

Using the Performance Profiler

- It's one of the tabs in DevTools
- Switch to that tab and press the record button while your program is running
- Instantly receive a LOT of information about what parts were slowest
 - Top-down breakdown: Drill down into functions based on percentage of total runtime
- Another button can capture slowness during initial load (it records the page refresh)





Preventing slowdown in graphics: Spatial Data Structures

First problem:

- Problem #1: You have (large #) of points. Find all pairs that are close together.
 - This is a search problem
 - What are some CS32 data structures for fast lookup?




Hash tables: "Spatial Hashing"

- We want some mapping of points to hash buckets:
 $f(x,y,z) \Rightarrow 0 \text{ through } 999$
(supposing 1000 buckets)
- Normally hash functions should be random
 - We'd lose neighborhoods that way
 - What we need is called "locality-sensitive hashing"
- Naive answer: Take floor of x,y,z then do:
 $f(x,y,z) = x + 10y + 100z.$



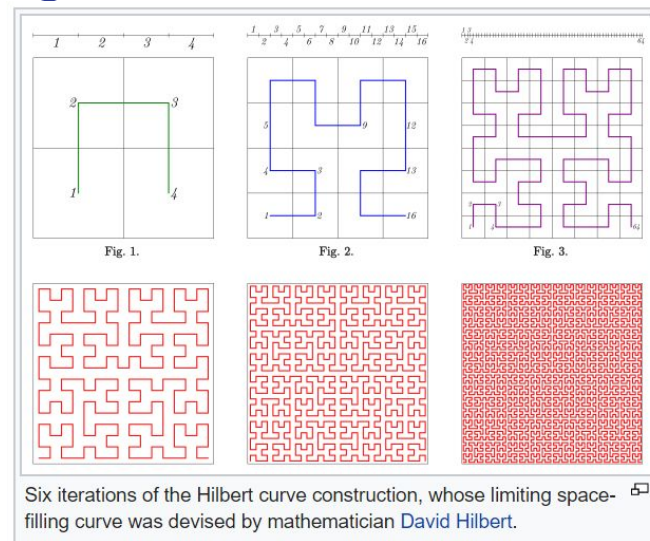
Hash tables: "Spatial Hashing"

- Naive answer: Take floor of x, y, z then do:
$$f(x, y, z) = x + 10y + 100z.$$
- Not cache-friendly; nearby objects one row away become far away in memory
- Better orderings: "Space-filling curve" order
 - Peano Curve: 1D curves that pass through every point in the unit square if you keep following them



Hash tables: "Spatial Hashing"

- Another one: Hilbert Curve



- Can generate a nice curve with "Morton Coding" (alternate the bits of the integers)

Hash tables: "Spatial Hashing"

- For now suppose any ordering, such as $f(x,y,z) = x + 10y + 100z$
 - Divides all of space into a 3D grid along integers
 - Does our scene have to be stuck in a 10x10x10 aquarium?
 - Let's try using x,y,z out of range anyway
 - Hopefully aliasing grid cells to same buckets results in an even distribution



Next Problem:

- Problem #2: You have (large #) matrices, each of "bounding boxes" for objects we want to draw. Find all pairs that are close together.
- Naive solution: Do the same thing as for points
- Problem: Now our points have size
 - Some objects will dip over a grid cell's edge
 - Some objects will completely mismatch our grid scale



Problem #2

- Some objects will dip over a grid cell's edge
- Some objects will completely mismatch our grid scale
- Solution 1: Be repetitive.
 - For each matrix, store a pointer to the object in *every* hash bucket (grid cell) that the object overlaps.



Problem #2

- Problem: What if the object is huge and overlaps many or all of the grid cells?
 - We won't have reduced the search space at all
 - Objects that have movement / velocity become expensive
 - Remove every grid cell with the matching pointer (recursively scanning in all directions), add to new cells



Problem #2

- Problem: What if the object is huge and overlaps many or all of the grid cells?
 - We won't have reduced the search space at all
 - Objects that have movement / velocity become very slow and expensive
 - Remove every grid cell with the matching pointer (recursively scanning in all directions), add to new cells



Problem #2

- Solution 2:
 - Prevent mis-matched object scale and grid scale by maintaining a lot of grids.
 - Each grid is 2x the size of the previous
 - A bounding box is stored in the smallest grid where at most it overlaps 8 cells (in case of sitting on edges)



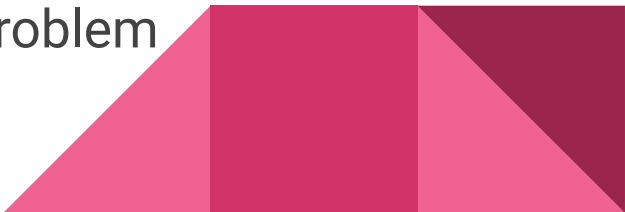
Problem #2

- Implementation: Multiple hash tables, each with maybe 1000 hash buckets
 - An object is represented in one table, up to 8 times via pointers
- Software using this technique: "VDB" package of Houdini (Fluid simulations)
 - Video of someone using the plugin for tracking turbulence details within a smoke simulation:

<https://www.youtube.com/watch?v=MR8CA6BaVm8>



Benefits of our Spatial Hash Table

- “Frustum culling”
 - Reduces the # of draw calls
 - Easier to search for only the objects that are in view
 - Collision detection
 - Now we reduced the problem from $O(n^2)$ to $O(1)$
 - With a fairly small constant too, due to 8 neighbors max
 - Ray tracing
 - We also reduced the ray / object collision problem down from $O(n \cdot \text{pixels} \cdot \text{recursions})$
- 

Benefits of our Spatial Hash Table

- Ray tracing
 - We also reduced the ray / object collision problem down from $O(n * \text{pixels} * \text{recursions})$
 - We still have to "march" the ray through the grid to the first struck object
 - Start at the biggest grid, have those cells pre-marked for if we need to drill down to a smaller grid
 - To march the ray (S, C) , solve for t along the ray $(S + tC)$ that first intersects a nearby grid plane (integers)



Alternatives

- Tetrahedralized models
 - Instead of representing 2D surfaces (shells) of every object, fill in and connect the interior volume too
 - Make every object out of the simplest 3D shape, a tetrahedron
 - Common in elastics and fluids simulation (where forces are volumetric)



Alternatives

- Tetrahedralized models
 - Main application for helping with spatial search:
 - **Tetrahedralize the air or empty space in between objects too**
 - Now everything is connected by edges
 - Edge connections always exist to tell you when objects about to collide
 - Ray tracing now just requires marching the ray across faces of neighboring tetrahedrons



Alternatives

- Example software: "Deformable Simplicial Complex"
 - Maintains this "mesh" even as objects move, by deforming tetrahedrons, adding & deleting when necessary
- Video 1:
 - My attempt (2D): <https://www.youtube.com/watch?v=STIRbYcVY6E>



Alternatives

- Video 2:
 - Their attempt (3D): <https://vimeo.com/48290450>
 - Tetras in the air are not drawn (invisible)
 - They can freely deform/appear/disappear
 - Tetras touching a surface are hard constraints bound to the sim's forces
 - Notice the triangles popping in and out of existence when the sim deforms it





What “real” WebGL programs can
do

“BioDigital” Web Human Anatomy Explorer

- Giant database of demos that run online and teach anatomy using WebGL models
- Knee example:
 - https://human.biodigital.com/widget/?m=production/maleAdult/musculoskeletal_system_knee.json



“After the Flood” WebGL Demo

- Made with PlayCanvas Engine
- <https://playcanv.as/e/p/44MRmJRU/>
- A virtual scene that pushes performance limits

